

Entity Resolution for Veridion Data Engineer Internship

A fast, reliable Approach to Noisy Company Records

(Bîrsan Gheorghe-Daniel)

Contents

1	Introduction	2
1.1	Context and Rationale	2
1.2	Motivation	2
1.3	Overall Two-Part Pipeline	2
2	Challenge & Constraints	3
2.1	The Task	3
2.2	Goals & Requirements	3
3	Solution Outline	3
3.1	Part 1: Data Normalization and Preprocessing	3
3.2	Part 2: Blocking and Finding Duplicates	4
3.3	Part 3: Clustering & Enriching the Data	4
4	Problems Encountered	4
4.1	Naive Approach with Column-by-Column Compare	4
4.2	Overly Detailed Location Enrichment	5
4.3	General Speed vs. Coverage Trade-off	5
5	Detailed Documentation of the Code	5
5.1	Data Normalization (Step 1)	5
5.1.1	Logic Explanation	5
5.1.2	Data Example	6
5.2	Finding Duplicates (Step 2)	6
5.2.1	Short Description of the Approach	6
5.2.2	Scoring System & Code Explanation	6
5.3	Clustering and Enriching Data (Step 3)	9
5.3.1	Short Description of the Approach	9
5.3.2	Merging Code (Detailed)	9
6	Analysis of Criteria	9
6.1	Correctness	9
6.2	Robustness	9
6.3	Code Quality	10
6.4	Extra Mile	10

6.5	Presentation	10
7	Future Adjustments	10
8	Conclusion	11
9	Results	11

1 Introduction

1.1 Context and Rationale

In modern data environments, companies often receive large, *noisy datasets* from multiple sources, each using different standards for naming, location, industry codes, and so on. The **task** here is to **identify unique companies** among these messy records, gather duplicates into clusters, and produce a *single enriched row* for each cluster. This challenge was presented specifically for a *Veridion Data Engineer Internship* assignment but is relevant to many data integration scenarios.

1.2 Motivation

- Data duplication leads to inflated statistics, confusion, and wasted resources.
- Combining partial or inconsistent data into the “best” or “golden” record yields the most robust dataset.
- Speed and correctness are both crucial: blocking or indexing ensures we do not blow up computationally.

1.3 Overall Two-Part Pipeline

In broad terms, the approach can be split into two phases:

1. Normalization & Preprocessing:

- Cleaning textual columns, normalizing websites, approximating missing location fields, merging or inferring `company_name` from other columns.
- Producing a `df_cleaned` that is consistent and more easily comparable.

2. Similarity & Clustering:

- Using a blocking strategy (e.g., `main_country`) to keep it fast.
- A scoring system (points-based or fuzzy approach) to decide if pairs are duplicates.
- Converting these duplicates into connected components and merging each cluster into a single row.

2 Challenge & Constraints

2.1 The Task

We are given a Parquet file with company records. The columns span:

- **Company Identifiers:** `company_name`, `company_legal_names`, ...
- **Location Data:** `main_country_code`, `main_region`, `main_city`, ...
- **Contact:** `primary_phone`, `phone_numbers`, ...
- **Web Presence:** `website_url`, `website_domain`, `facebook_url`, ...
- **Business Details:** `company_type`, `year_founded`, ...
- **Industry Codes:** `naics_vertical`, `main_business_category`, ...
- **Other Data:** `revenue`, `employee_count`, `inbound_links_count`, ...

The dataset is known to be:

- *Noisy* (missing, partial, or contradictory fields).
- *Inconsistent* (some records only have `company_name`, others only have a partial `company_legal_names`, etc.).
- Potentially *large*, requiring efficient grouping of duplicates.

2.2 Goals & Requirements

1. **Speed:** Must handle tens (or hundreds, I hope not :*(but it is quite fast) of thousands of rows quickly.
2. **Correctness:** Identify duplicates with minimal false positives or false negatives.
3. **Robustness:** Should handle unexpected or missing data gracefully.
4. **Development Feasibility:** Code should be maintainable, understandable, and easily adapted for future improvements.

3 Solution Outline

The final approach divides the solution into **three major parts**:

3.1 Part 1: Data Normalization and Preprocessing

- **Enrich Company Name:** If `company_name` is missing, use `company_legal_names`, or `short_description`, or parse the `website_domain`.
- **Location Approximation:** If `main_country` is missing, parse from `main_country_code` or from TLD. We skip advanced geocoding to remain offline and fast.

- **Activity Sector Normalization:** Combine codes like `naics_2022_primary_code`, `sic_codes`, or parse textual columns with an ML model.
- **City and Country Finder:** Extract City and Country from the Postcode.
- **Numeric Columns:** Convert `revenue`, `employee_count` to numeric if possible, ignoring invalid strings.

3.2 Part 2: Blocking and Finding Duplicates

- **Blocking by Country:** Only compare records within the same country. If `main_country` is missing, group those in a “*missing*” or “*unknown*” block.
- **Scoring System:** Assign points (or partial fuzzy scores) for matching columns:
 - **Company Name** (fuzzy ratio via RapidFuzz).
 - **Website Domain** (exact match).
 - **Social Media** links (partial or exact match).
 - **Revenue** or **Employee Count** if within $\pm 10\%$.
- **Thresholding:** If the total points exceed (e.g.) 4, mark them as duplicates.

3.3 Part 3: Clustering & Enriching the Data

- **Connected Components:** Build a graph where each record is a node, each duplicate pair is an edge. All records in the same connected component are duplicates of each other.
- **Merging Rows:** For each cluster, pick the best data:
 - *Longest textual field* (ensures maximum detail).
 - *Largest numeric* or *average*, depending on domain preference.
 - *First non-null* fallback if no others exist.
- **Save Output:** The final DataFrame of deduplicated rows is stored in `final_data.parquet`.

4 Problems Encountered

Despite the plan, we ran into:

4.1 Naive Approach with Column-by-Column Compare

A first naive approach attempted to compare every row with every other row and check all columns if not null. This took **13 hours** and was obviously too slow.

4.2 Overly Detailed Location Enrichment

We tried geocoding via external APIs, but:

- Incurred network latency.
- Data usage / external calls overhead.
- Not feasible for a purely offline, time-sensitive solution.

Hence, we pivoted to TLD-based or code-based approximations and used a database found online which we used to extract city and country from the postcode.

4.3 General Speed vs. Coverage Trade-off

While more advanced or multi-level blocking might increase coverage, I had to ensure I didn't lose reliability. Ultimately, a single country-based block plus an optional fallback for missing countries was sufficient. I also tried web scrapping, to ensure the data but it was a time consuming running process so I gave up, because of hardware and time constraints. Also, it was possible to enrich more the activity codes, but I didn't succeed so I tried to focus on other criterias giving the development time constraint

5 Detailed Documentation of the Code

5.1 Data Normalization (Step 1)

5.1.1 Logic Explanation

- `best_company_name(row)` obtains the best possible name by:
 1. Checking `company_name` or fallback to `company_legal_names`.
 2. If missing, parse `short_description` or `generated_description`.
 3. If still missing, parse `website_domain` (remove suffixes like "inc" or "ltd").
- `fill_missing_location(row)` tries to approximate `main_country` from `main_country_code` or from the TLD of `website_url`.
- `unify_activity_info(row)` merges multiple industry codes and short descriptions into a single `activity_enriched` field.

Illustration Code (Excerpt):

```
def best_company_name(row: dict) -> str:
    web = row.get("website_url")
    if pd.isna(web):
        web = None

    if web:
        d, _ = parse_domain_tld(web)
        if d:
            return remove_company_suffixes(normalize_text(d))

    # Fallback logic with short_description, etc.
```

```

for desc_col in ["short_description", "long_description", "generated_description"]:
    val = row.get(desc_col)
    if val and isinstance(val, str):
        # use first n words
        first3 = get_first_n_words(val, 3)
        return remove_company_suffixes(normalize_text(first3)) or None

# final fallback: company_name or commercial_names
...
return None

```

5.1.2 Data Example

Imagine a row with:

- `website_url` = `http://www.testcompany.com`
- `company_name` = `NA`
- `short_description` = “leading test solutions”

I parse out “testcompany” from `website_url` if no better name is found.

5.2 Finding Duplicates (Step 2)

5.2.1 Short Description of the Approach

We block by `main_country`. So it do not compare records from the US with records from Germany (unless the country is missing). We define a function `compute_pair_score(rowA, rowB)` awarding points for:

- **Name similarity** (via `rapidfuzz.fuzz.ratio`).
- **Exact domain match** (some columns are strictly 1 point or 2 points if identical).
- **Revenue or Employee Count** within $\pm 10\%$.
- **Social media** links if exact matches.

If total points exceed a threshold (like 4), we mark them as a likely duplicate pair.

5.2.2 Scoring System & Code Explanation

```

def within_10pct(valA, valB):
    if valA is None or valB is None:
        return False
    try:
        vA = float(valA); vB = float(valB)
    except:
        return False
    if vA == 0 or vB == 0:
        return False
    ratio = vA / vB
    return (0.9 <= ratio <= 1.1)

def compute_pair_score(rowA, rowB):

```

```

points = 0
# Name fuzzy => up to 2 points
ratio = fuzz.ratio(rowA["company_name"] or "", rowB["company_name"] or "")
if ratio >= 80: points += 2
elif ratio >= 50: points += 1

# Website domain exact => +2
if rowA["website_domain"] and rowA["website_domain"] == rowB["website_domain"]:
    points += 2

# Revenue or employee_count within +/- 10% => +1 each
if within_10pct(rowA["revenue"], rowB["revenue"]):
    points += 1
if within_10pct(rowA["employee_count"], rowB["employee_count"]):
    points += 1

# Bonus for social links
# ...
# Punishment for different cities/regions
# ...
return points

```

Similarity Scoring Function

The similarity score $S(A, B)$ between two company records A and B is calculated as a weighted sum of multiple features:

$$\begin{aligned}
S(A, B) = & \text{NameScore}(A, B) \\
& + 0.5 \cdot 1_{\text{revenue within } \pm 30\%} \\
& + 0.5 \cdot 1_{\text{employee_count within } \pm 30\%} \\
& + 2 \cdot 1_{\text{website_domain match}} \\
& + \sum_{i=1}^n 1_{\text{social}_i(A)=\text{social}_i(B)} \\
& + \text{CityScore}(A, B) \\
& + \text{CountryScore}(A, B) \\
& + \text{RegionScore}(A, B) \\
& + \text{ActivitySim}(A, B)
\end{aligned}$$

Name Similarity (NameScore) based on fuzzy string matching:

$$\text{NameScore}(A, B) = \begin{cases} +1.5 & \text{if } \text{fuzz}(A, B) \geq 80 \\ +1.0 & \text{if } 50 \leq \text{fuzz}(A, B) < 80 \\ +0.5 & \text{if } 30 \leq \text{fuzz}(A, B) < 50 \\ -1.0 & \text{if } 10 \leq \text{fuzz}(A, B) < 30 \\ -3.0 & \text{if } \text{fuzz}(A, B) < 10 \end{cases}$$

City Similarity (CityScore):

$$\text{CityScore}(A, B) = \begin{cases} -0.2 & \text{if cities differ and no contextual match} \\ +0.1 & \text{if city found in location/address context} \\ 0 & \text{if cities match or missing} \end{cases}$$

Country Similarity (CountryScore):

$$\text{CountryScore}(A, B) = \begin{cases} -2.0 & \text{if countries differ and no contextual match} \\ +0.2 & \text{if country appears in contextual location fields} \\ 0 & \text{if countries match or missing} \end{cases}$$

Region Similarity (RegionScore):

$$\text{RegionScore}(A, B) = \begin{cases} -1.0 & \text{if regions differ and not in context} \\ +0.1 & \text{if region match is found in location fields} \\ 0 & \text{if regions match or missing} \end{cases}$$

Activity Similarity (ActivitySim): Based on Jaccard index J over the sets of extracted activities:

$$J = \frac{|\text{activities}(A) \cap \text{activities}(B)|}{|\text{activities}(A) \cup \text{activities}(B)|}$$

$$\text{ActivitySim}(A, B) = \begin{cases} +0.5 & \text{if } J \geq 0.7 \\ +0.2 & \text{if } 0.5 \leq J < 0.7 \\ -0.5 & \text{if } J < 0.5 \\ 0 & \text{if missing} \end{cases}$$

Legend

- $1_{\text{condition}}$: Indicator function — 1 if condition is true, 0 otherwise.
- $\text{fuzz}(A, B)$: Fuzzy string matching score (0 to 100) between company names.
- website_domain : Exact domain match contributes +2 points.
- social_i : Social media fields (e.g., Facebook, LinkedIn) contribute +1 point per match.
- $\text{within } \pm 30\%$: For numeric fields like revenue or employee count.
- contextual match : Checks if mismatched location fields appear in raw text or other known location columns.

Final score $S(A, B)$ is lower-bounded at 0:

$$S(A, B) = \max(0, \text{sum of all components})$$

After computing these pairwise scores, we store pairs with `score >= 4` in `df_duplicates`.

5.3 Clustering and Enriching Data (Step 3)

5.3.1 Short Description of the Approach

We treat each record as a graph node. If two records are duplicates, we add an edge. Then **all connected nodes** belong to the same cluster. We unify them with a function that picks the best column value:

- *Longest text* for textual columns.
- *Largest or first non-null* for numeric columns.

5.3.2 Merging Code (Detailed)

```
import networkx as nx

G = nx.Graph()
for row in df_duplicates.itertuples():
    idxA = row.idxA
    idxB = row.idxB
    G.add_edge(idxA, idxB)

components = list(nx.connected_components(G))

def merge_cluster(cluster_indices, df):
    subset = df.loc[list(cluster_indices)]
    merged_row = {}
    for col in df.columns:
        colvals = subset[col].dropna().tolist()
        if len(colvals) == 0:
            merged_row[col] = None
            continue
        # if text, pick longest
        # if numeric, pick largest, etc.
    return merged_row
```

We then iterate through each connected component, *merge* those rows, and produce a final DataFrame. Finally, we write that DataFrame to `final_data.parquet`.

6 Analysis of Criteria

6.1 Correctness

- By carefully normalizing columns and awarding points for name, domain, location, etc., we reduce the chance of false positives.
- Using *connected components* ensures that if A matches B, and B matches C, then A, B, C are all considered duplicates.

6.2 Robustness

- Missing columns default to `None`, skipping or awarding 0 points.
- We do not rely on external APIs for location; thus no blocking network calls.
- If the dataset is partially large, we can scale with Dask or Spark if needed.

6.3 Code Quality

- Each transformation is in a clear function: `best_company_name`, `fill_missing_location`, ...
- The final merging code is also modular.
- Comments and docstrings are used for clarity.

6.4 Extra Mile

- We used **longest text** logic in merging to preserve maximum detail.
- We handled optional columns like social media and partial phone matches for partial scoring.
- We gave expansions for *revenue* or *employee_count* within $\pm 10\%$.

6.5 Presentation

We carefully outlined the approach in steps, each with code snippets and logic explanation. We also included a verification snippet that prints original cluster rows for a quick sanity check.

7 Future Adjustments

Even after building a stable approach, we foresee potential expansions:

- **LSH (Locality-Sensitive Hashing)**: If the name fields are extremely large or we approach millions of records, an LSH-based approach can drastically reduce pairwise comparisons.
- **Weighted Merging**:
 - Instead of simply “largest numeric” or “longest text,” we can track *confidence* in each record or do a majority vote.
- **Language Model Summaries**: We might unify textual columns by concatenating or summarizing them with an ML model to produce the final text description.
- **Advanced k -Means or DBSCAN Over Embeddings**: Instead of pairwise scoring, embed the records in a vector space (like a *Transformer-based* embedding) and use clustering in high-dimensional space. This might be slower or more complex but could capture deeper semantic similarities.
- **Spark-based Graph Merging**: If the data is extremely large, we can use Apache Spark GraphX or GraphFrames for distributed connected-component detection.

8 Conclusion

1. **We split the problem** into a first step for *data normalization* and a second step for *finding and merging duplicates*.
2. **Normalization** ensures each column (e.g. `company_name`, `main_country`, `revenue`) is consistent, numeric fields are typed properly, and textual fields are trimmed.
3. **Similarity** uses blocking on country, a points-based system for columns, and a threshold for duplicates.
4. **Clustering** is done by turning duplicates into a *graph* and identifying *connected components*.
5. **Merging** finalizes the golden record, choosing the best textual or numeric values.

Thus, the final DataFrame in `final_data.parquet` is *fully deduplicated* and *enriched* with the best available data from each cluster of duplicates.

Why this Approach?

- It is *fast* enough for moderate to large data, especially with an offline, TLD-based location fill and a country-based block.
- It is *robust* in that we handle missing data gracefully.
- It is *easy to adjust* the scoring system or the merging logic as domain knowledge evolves.

9 Results

After executing the full deduplication pipeline on the cleaned dataset, the following outcomes were obtained:

- **Total Connected Components (Clusters):** 19,779
 - A component with only one node represents a unique company without detected duplicates.
 - Components with more than one node indicate detected duplicates, grouped together.
- **Potential Duplicate Pairs Found (score \geq 4):** 23895 ‘
 - These represent candidate record pairs with high similarity based on the scoring system.
 - The score threshold of 4 was determined based on empirical testing to balance recall and precision.
 - The scoring system included fuzzy name match, website match, employee count similarity, revenue range, and social media overlap.

These results show that nearly a third of the dataset involved potential duplication, underlining the importance of this deduplication process. The graph-based clustering then ensured all linked duplicate pairs were unified into single golden records.

The final deduplicated and enriched dataset was saved as:

- **Filename:** `final_data.parquet`
- **Format:** Apache Parquet (columnar storage, optimized for large-scale analytics)

Final Words: The *Entity Deduplication* method proposed meets the requirements for a **Veridion Data Engineer Internship** by balancing speed, correctness, running time, and development effort. Our detailed code and approach demonstrate professional best practices in data normalization, fuzzy matching, cluster merging, and overall pipeline structure.